

Minemu: Protecting buggy binaries from memory corruption attacks

Programming Languages

type-safe

vs.

not type-safe

Programming Languages

type-safe

vs.

not type-safe

Java

Python

Ruby

Javascript

Programming Languages

type-safe

vs.

not type-safe

Java

C

Python

C++

Ruby

Javascript

Programming Languages

type-safe

vs.

not type-safe

Java

Python

Ruby

Javascript

C

C++

**MEMORY
CORRUPTIONS!**

Programming Languages

type-safe

vs.

not type-safe

Java

Python

Ruby

Javascript

**MEMORY
CORRUPTIONS!**

C

C++

**MEMORY
CORRUPTIONS!**

Programming Languages

type-safe

vs.

not type-safe

Java

Python

Ruby

Javascript

MEMORY

CORRUPTIONS!

but not

your fault

C

C++

MEMORY

CORRUPTIONS!

The Stack

[code]

```
run(char *name)
{
    char buf[16];

    print("hello ");
    print("world\n")
}
```


The Stack

[code]

```
run(char *name)
{
    char buf[16];

    print("hello ");
    print("world\n")
}
```

[stack]

b	a	s	e	r	e	t	n	a	r	g	1	.	.	.
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

The Stack

[code]

```
run(char *name)
{
    char buf[16];

    print("hello ");
    print("world\n")
}
```

[stack]

base ret narg 1 ...

buf base ret narg 1 ...

The Stack

[address] [code]

[stack]

```
8048751: run(char *name)
        {
```

```
        char buf[16];
```

```
8048770:     print("hello ");
```

```
8048798:     print("world\n");
```

```
}
```

baseretnarg1...

buf 16 bytes of memory (red) followed by baseretnarg1...

The Stack

[address] [code]

[stack]

```
8048751: run(char *name)
      {
```

```
      char buf[16];
```

```
8048770: print("hello ");
```

```
8048798: print("world\n")
```

```
}
```

base ret narg1 ...

buf base ret narg1 ...

ret narg1 buf base ret narg1 ...

The Stack

[address] [code]

[stack]

```
8048751: run(char *name)
      {
```

```
      char buf[16];
```

```
8048770: print("hello ");
```

```
8048798: print("world\n")
```

```
}
```

baseretnarg1...

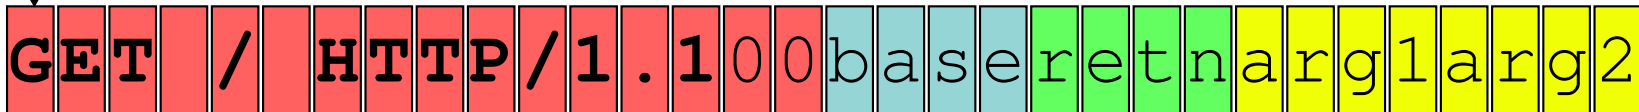
buf

retnarg1buf

retnarg1buf

Traditional Stack Smashing

buf[16]



Traditional Stack Smashing

buf[16]



GET / HTTP/1.100base ret narg1arg2

SHELLCODE!@#\$%^&*()_&buf

Address Space Layout Randomisation (ASLR)

buf[16]

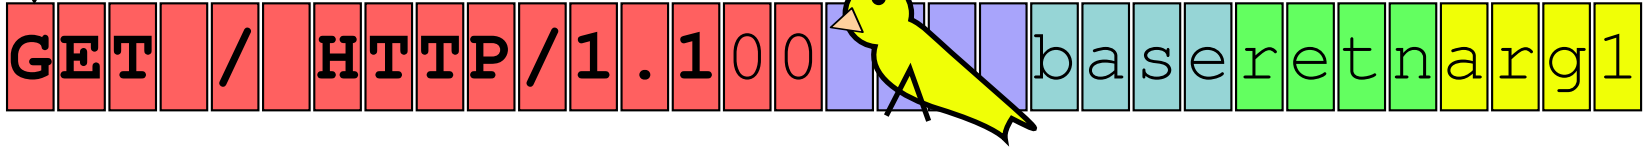


G	E	T		/		H	T	T	P	/	1	.	1	0	0	b	a	s	e	r	e	t	n	a	r	g	1	a	r	g	2
---	---	---	--	---	--	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

S	H	E	L	L	C	O	D	E	!	@	#	\$	%	^	&	*	()	_	?	?	?	?								
---	---	---	---	---	---	---	---	---	---	---	---	----	---	---	---	---	---	---	---	---	---	---	---	--	--	--	--	--	--	--	--

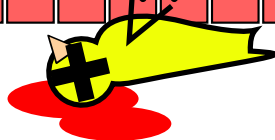
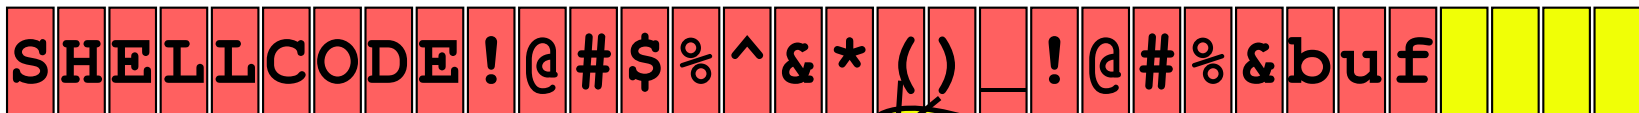
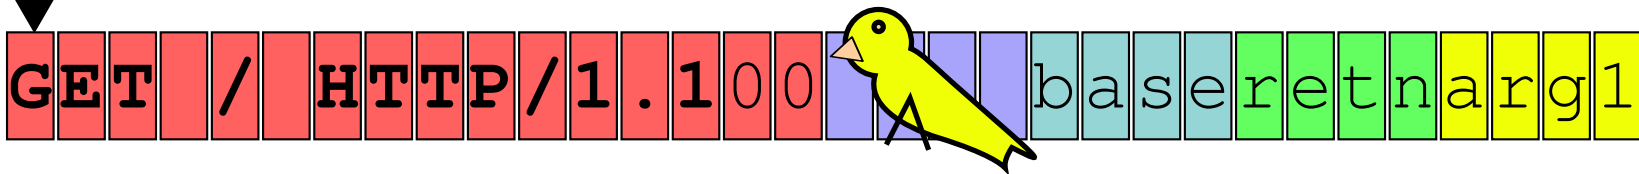
Stack Canaries

buf[16]



Stack Canaries

buf[16]



Non-executable data (DEP / NX)

buf[16]



GET		/		HTTP	/	1	.	1	0	0	b	a	s	e	r	e	t	n	a	r	g	1	a	r	g	2
-----	--	---	--	------	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

S	H	E	L	C	O	D	E	!	@	#	\$	%	^	&	*	()	_	&	b	u	f								
--------------	--------------	--------------	--------------	--------------	--------------	--------------	--------------	---	---	---	----	---	---	---	---	---	---	---	---	---	---	---	--	--	--	--	--	--	--	--

Fortify Source

```
char buf[16];
memcpy(buf, r->buf, r->len);
```

GET / HTTP/1.1baseretnarg1arg2

[illegible]

Fortify Source

```
char buf[16];  
memcpy(buf, r->buf, r->len);
```

G	E	T		/		H	T	T	P	/	1	.	1	0	0	b	a	s	e	r	e	t	n	a	r	g	1	a	r	g	2
---	---	---	--	---	--	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

```
char buf[16];  
memcpy_chk(buf, r->buf, r->len, 16);
```

s	h	;	S	T	A	C	K	S	M	A	S	H	E	R	A	A	A	A	A	A	A	A	A	A	A	A	A	A	A	A	A	A
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---



*** buffer overflow detected ***: /my/fortified/binary terminated

===== Backtrace: =====

/lib/i386-linux-gnu/i686/cmov/libc.so.6(__fortify_fail+0x50)[0xb774a4d0]

/lib/i386-linux-gnu/i686/cmov/libc.so.6(+0xe040a)[0xb774940a]

/my/fortified/binary[0x8048458]

/lib/i386-linux-gnu/i686/cmov/libc.so.6(__libc_start_main+0xe6)[0xb767fe46]

/my/fortified/binary[0x8048371]

===== Memory map: =====

08048000-08049000 r-xp 00000000 fe:00 282465

/my/fortified/binary

08049000-0804a000 rw-p 00000000 fe:00 282465

/my/fortified/binary

08600000-08621000 rw-p 00000000 00:00 0

[heap]

b764b000-b7667000 r-xp 00000000 fe:00 131602

/lib/i386-linux-gnu/libgcc_s.so.1

b7667000-b7668000 rw-p 0001b000 fe:00 131602

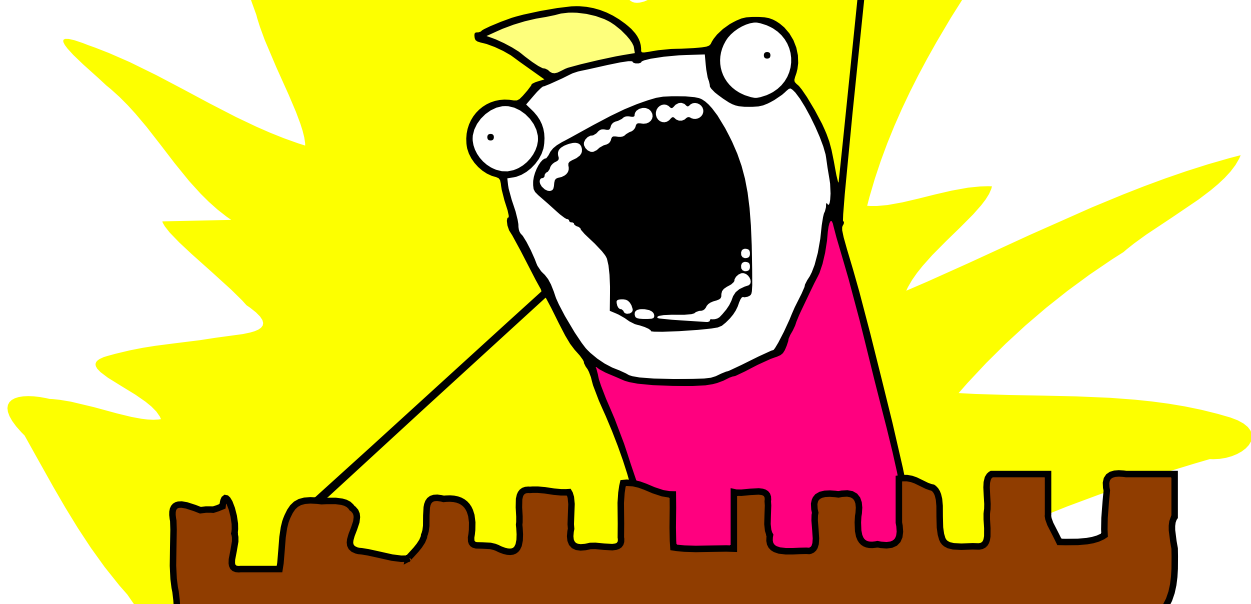
/lib/i386-linux-gnu/libgcc_s.so.1

b7668000-b7669000 rw-p 00000000 00:00 0

...

Aborted

FORTIFY SOURCE ALL THE THINGS!



Return Oriented Programming (ROP)

buf[16]



GET	/	HTTP	/	1	.	1	0	0	b	a	s	e	r	e	t	n	a	r	g	1	a	r	g	2
-----	---	------	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

sh	;	S	T	A	C	K	S	M	A	S	H	E	R	R	O	P	1	R	O	P	2	v	a	r	1
----	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

pointer to useful code



Some exploits still work with all these defense measures.

Example: nginx buffer underrun (CVE-2009-2629)

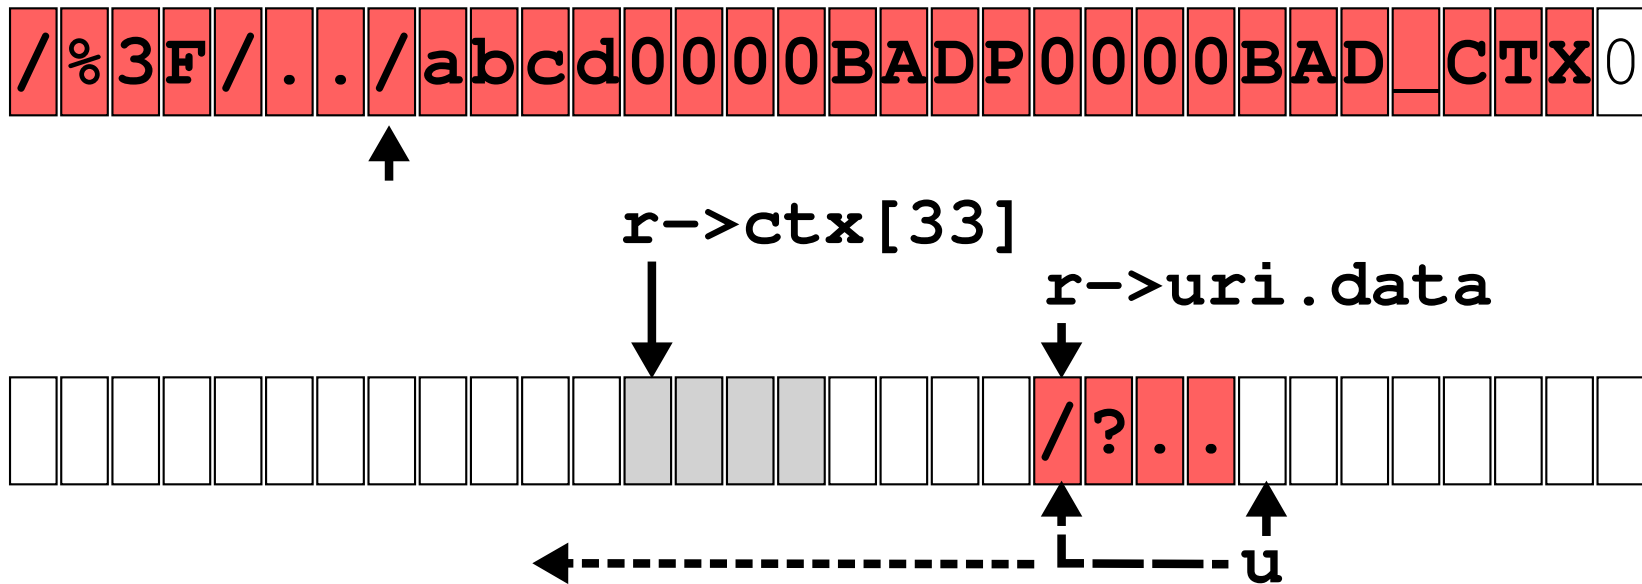
CVE-2009-2629

/	%	3	F	/	.	.	/	a	b	c	d	0	0	0	0	B	A	D	P	0	0	0	0	B	A	D	_	C	T	X	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

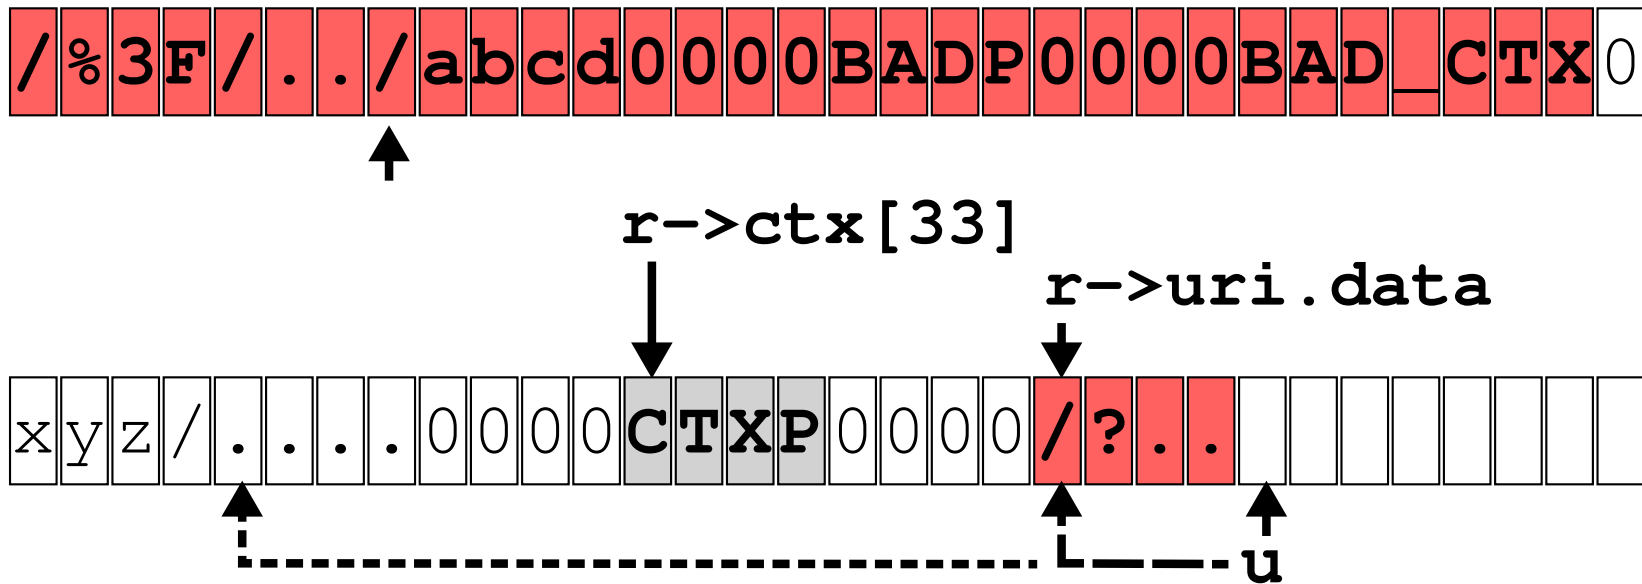


`r->uri_start`

CVE-2009-2629



CVE-2009-2629



CVE-2009-2629

/	%	3	F	/	.	.	/	a	b	c	d	0	0	0	0	B	A	D	P	0	0	0	0	B	A	D	_	C	T	X	0
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

`r->ctx[33]`

`r->uri.data`

x	y	z	/	0	0	0	0	B	A	D	P	0	0	0	0	B	A	D	_	C	T	X	0				
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	---	--	--	--	--



```
typedef struct {
    ngx_buf_t          *buf;
    ngx_chain_t        *in;
    ngx_chain_t        *free;
    ngx_chain_t        *busy;

    unsigned            sendfile;
    unsigned            need_in_memory;
    unsigned            need_in_temp;

    ngx_pool_t         *pool;
    ngx_int_t          allocated;
    ngx_bufs_t         bufs;
    ngx_buf_tag_t      tag;

    ngx_output_chain_filter_pt output_filter;
    void                *filter_ctx;
} ngx_output_chain_ctx_t;
```

```
typedef struct {  
    ngx_buf_t  
    ngx_chain_t  
    ngx_chain_t  
    ngx_chain_t  
  
    unsigned  
    unsigned  
    unsigned  
  
    ngx_pool_t  
    ngx_int_t  
    ngx_bufs_t  
    ngx_buf_tag_t  
  
    ngx_output_chain_filter_pt  
    void  
} ngx_output_chain_ctx_t;  
  
*buf;  
*in;  
*free;  
*busy;  
  
sendfile;  
need_in_memory;  
need_in_temp;  
  
*pool;  
allocated;  
bufs;  
tag;  
  
output_filter;  
*filter_ctx;
```

function pointer




```
805ba93:  mov    (%ecx),%ebx          ; copy filename
          movl   $0x3,0x10(%ecx)
          mov    %ecx,(%esp)
          call   *0x2c(%ecx)
```


```
805ba93:  mov    (%ecx),%ebx          ; copy filename
        movl  $0x3,0x10(%ecx)
        mov    %ecx, (%esp)
        call  *0x2c(%ecx)

      ┌──────────────────────────┐
8052267:  mov    %eax,0x4(%esp)       ; push argv
        mov    %ebx, (%esp)   ; push filename
        call  *0x14(%ebx)
```

```
805ba93:  mov    (%ecx),%ebx          ; copy filename
        movl  $0x3,0x10(%ecx)
        mov    %ecx, (%esp)
        call  *0x2c(%ecx)

8052267:  mov    %eax,0x4(%esp)       ; push argv
        mov    %ebx, (%esp)   ; push filename
        call  *0x14(%ebx)

804b274:  <execve@plt>                ; get shell
```



- **defeats address randomisation (through info leak)**

- defeats address randomisation (through info leak)
- defeats non-executable data protection

- defeats address randomisation (through info leak)
- defeats non-executable data protection
- no standard copy function (no fortify src protections)

- defeats address randomisation (through info leak)
- defeats non-executable data protection
- no standard copy function (no fortify src protections)
- not return oriented, so stack smash protection does not matter

But the situation is even worse

But the situation is even worse

- needs to be enabled at compile time, and there is a lot of old code out there

But the situation is even worse

- needs to be enabled at compile time, and there is a lot of old code out there
- many packages do not apply these defence mechanisms even today

But the situation is even worse

- needs to be enabled at compile time, and there is a lot of old code out there
- many packages do not apply these defence mechanisms even today
- implementation flaws

Can we do more?

Can we do more?

>> Non-executable data prevents untrusted data from
being run as code

Can we do more?

- >> Non-executable data prevents untrusted data from being run as code
- << Return oriented programming replaces untrusted code with pointers to original code.

Can we do more?

>> Non-executable data prevents untrusted data from being run as code

<< Return oriented programming replaces untrusted code with pointers to original code.

>> Can we prevent untrusted pointers from being used as jump addresses?

Taint analysis

[illegible]

Taint tracking (1/2):

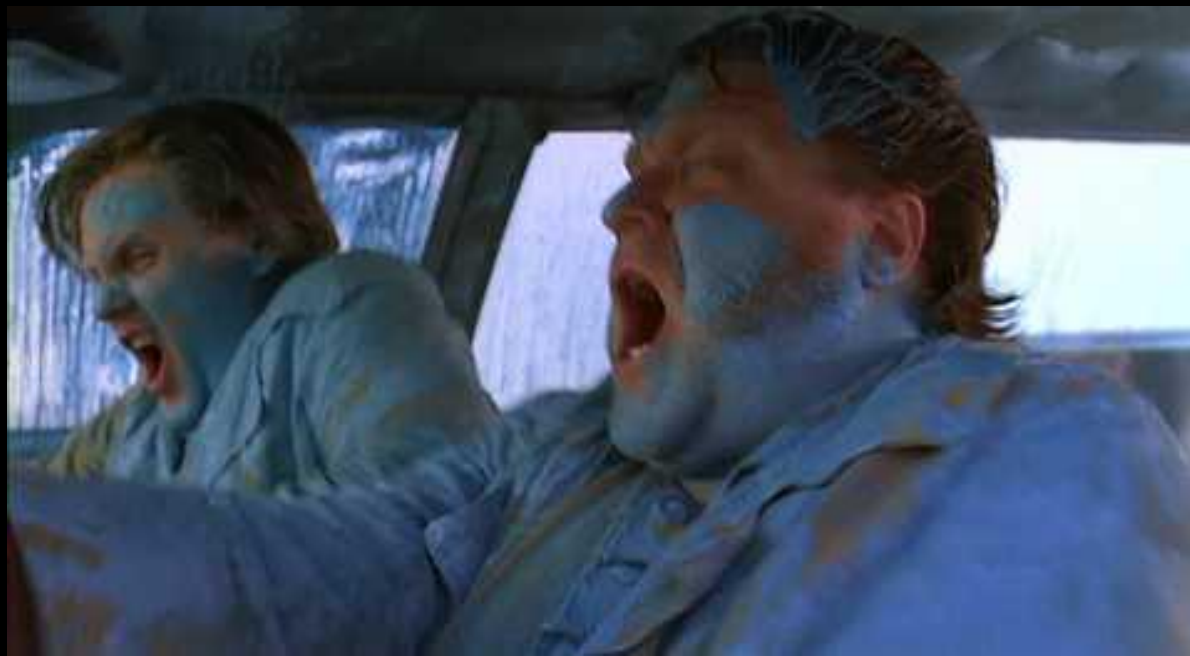
- remember whether data is trusted or not
- untrusted data is 'tainted'
- when data is copied, its taint is copied along
- taint is ORed for arithmetic operations

Taint tracking (2/2):

When the code jumps to an address in memory,
the source of this address is checked for taint.

eg.:

- RET
- CALL ***%eax**
- JMP ***0x1c(%ebx)**

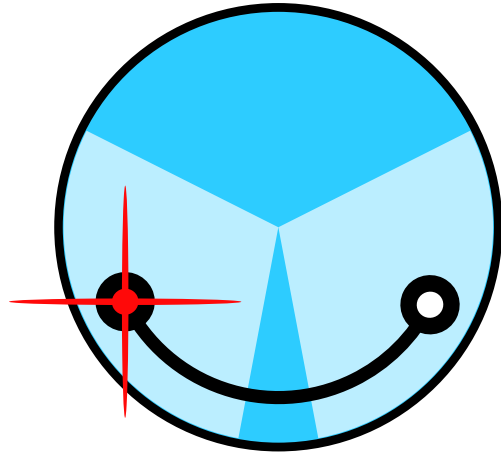


Taint tracking



useful, but slow as hell

Is this slowness fundamental?



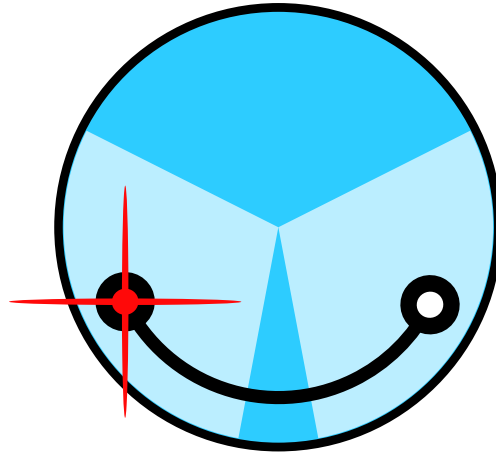
minemu

fast emulator

memory layout

use SSE registers to hold taint

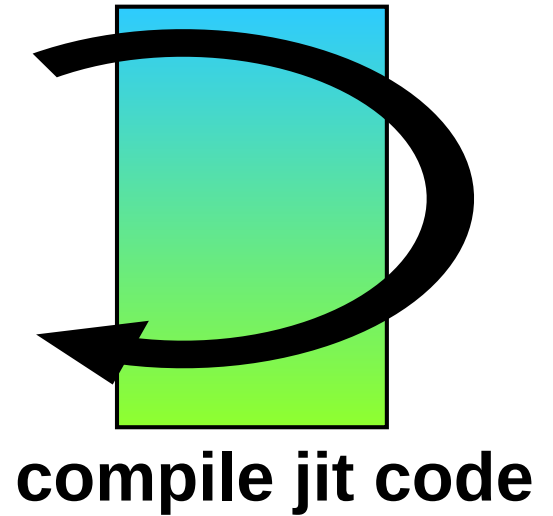
Is this slowness fundamental?



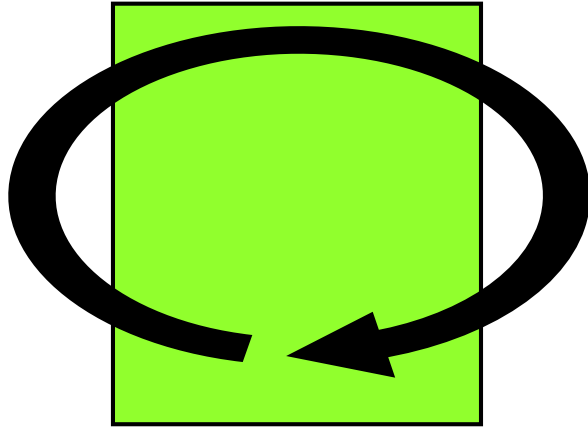
minemu

- ▶ fast emulator
- memory layout
- use SSE registers to hold taint

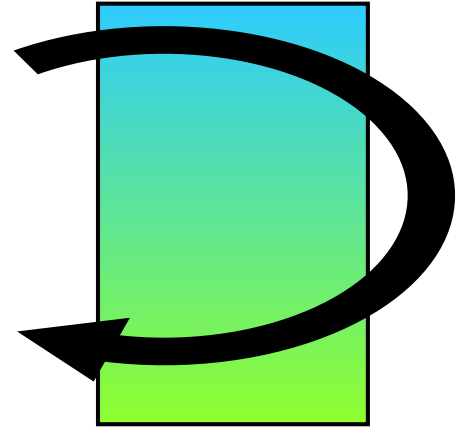
Emulator



Emulator

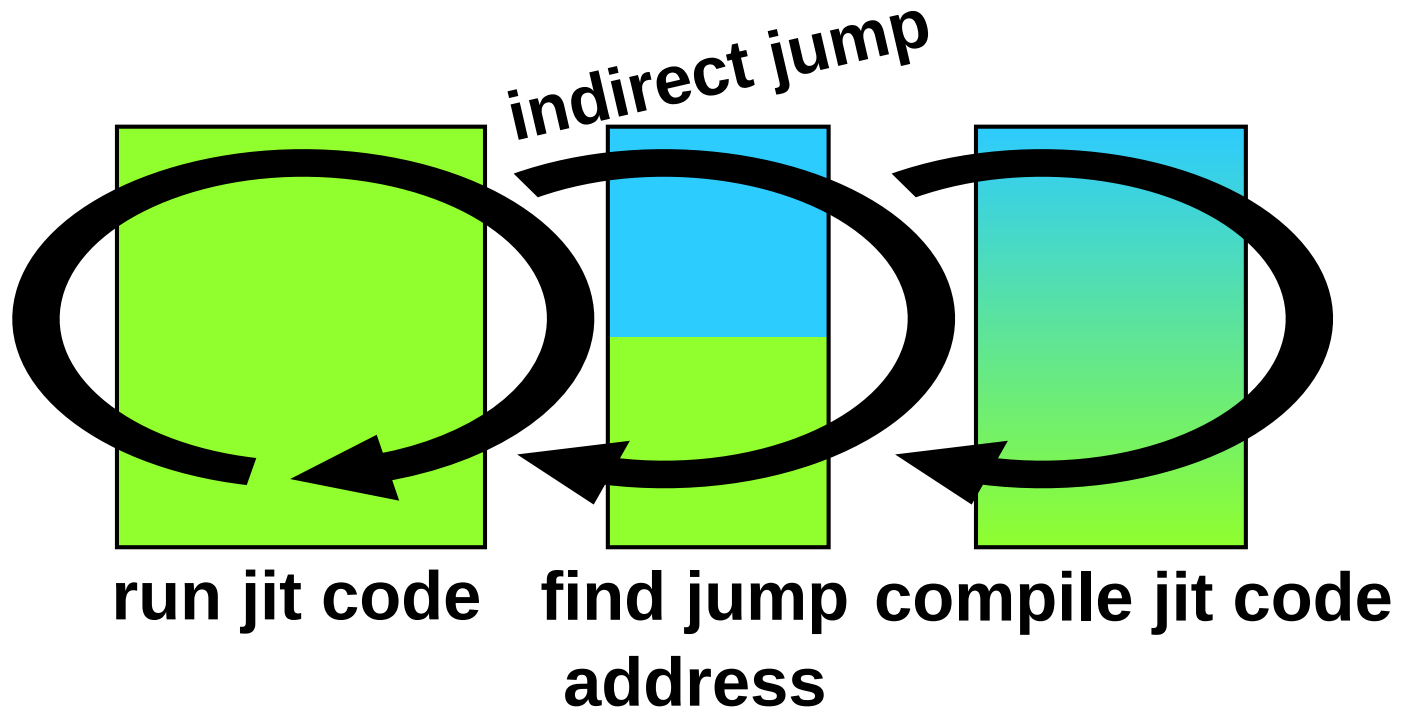


run jit code

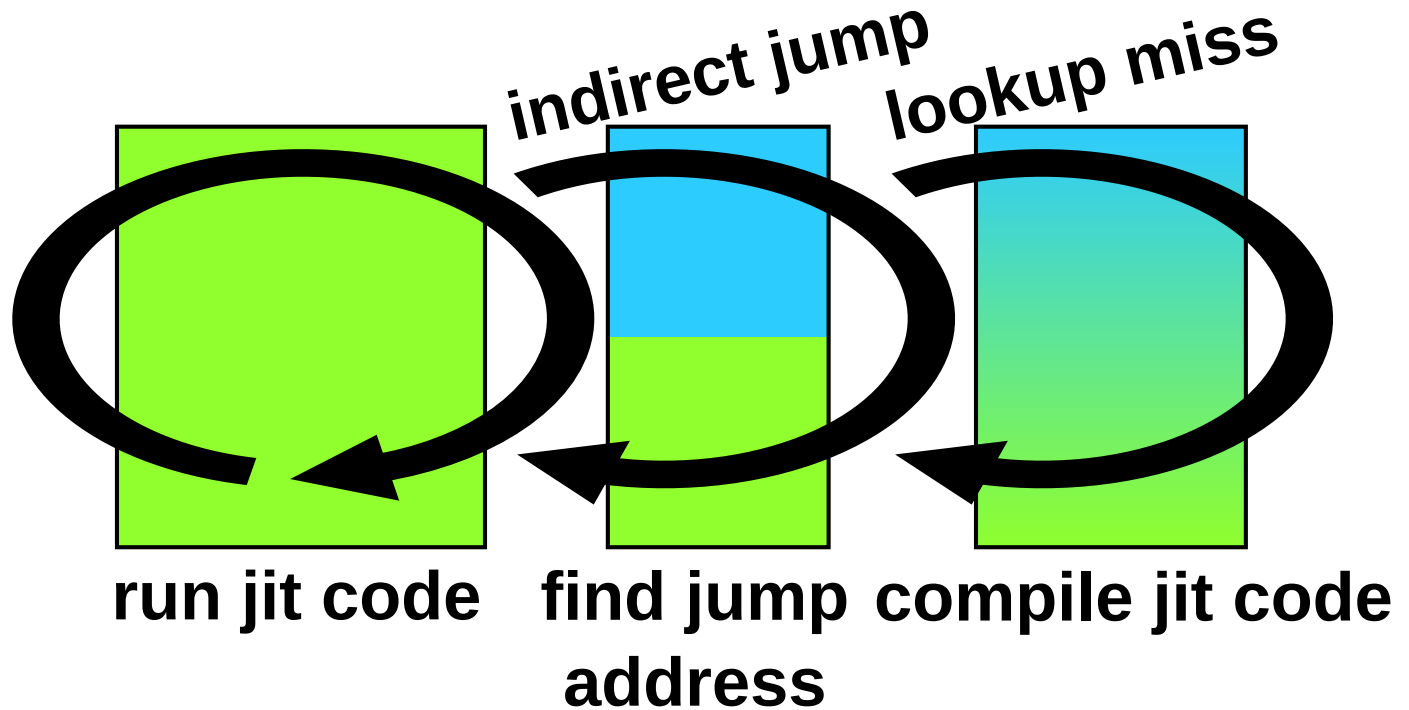


compile jit code

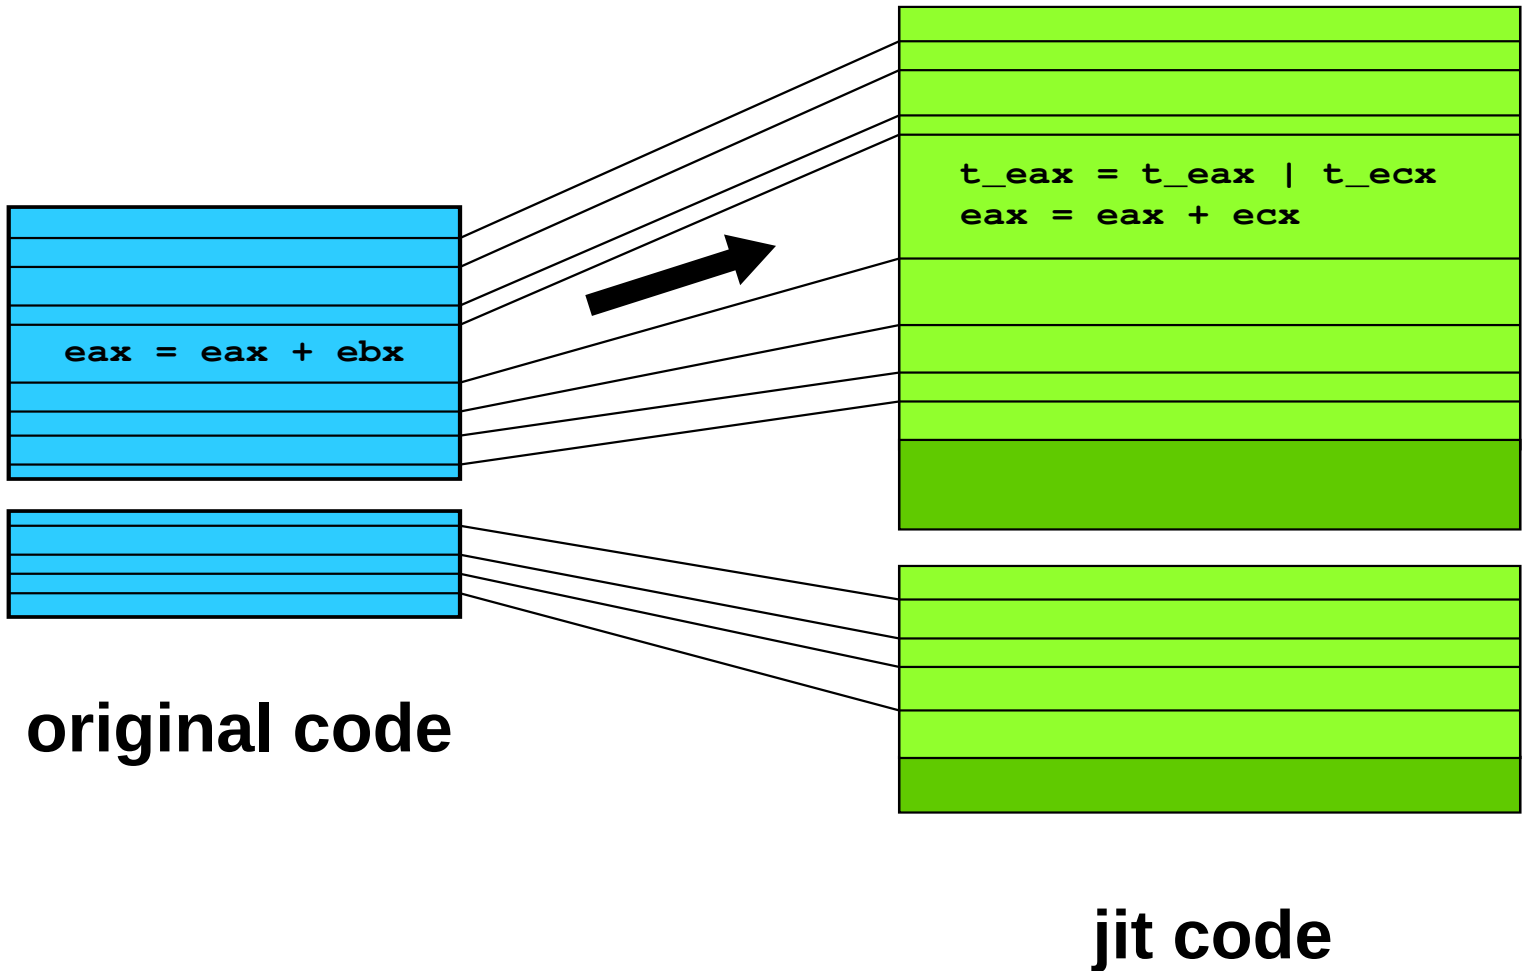
Emulator



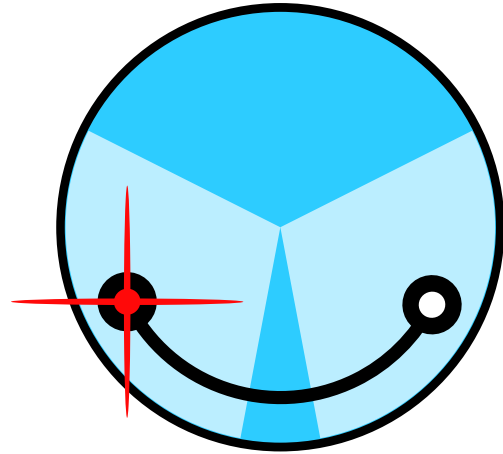
Emulator



Dynamic instrumentation



Is this slowness fundamental?



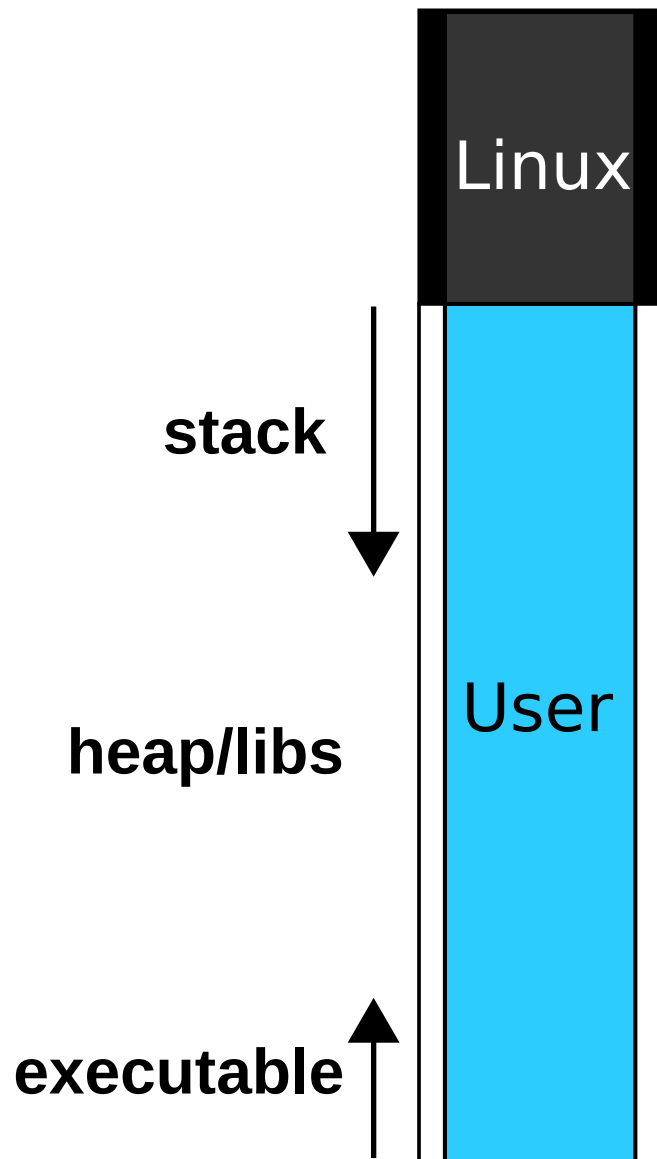
minemu

fast emulator

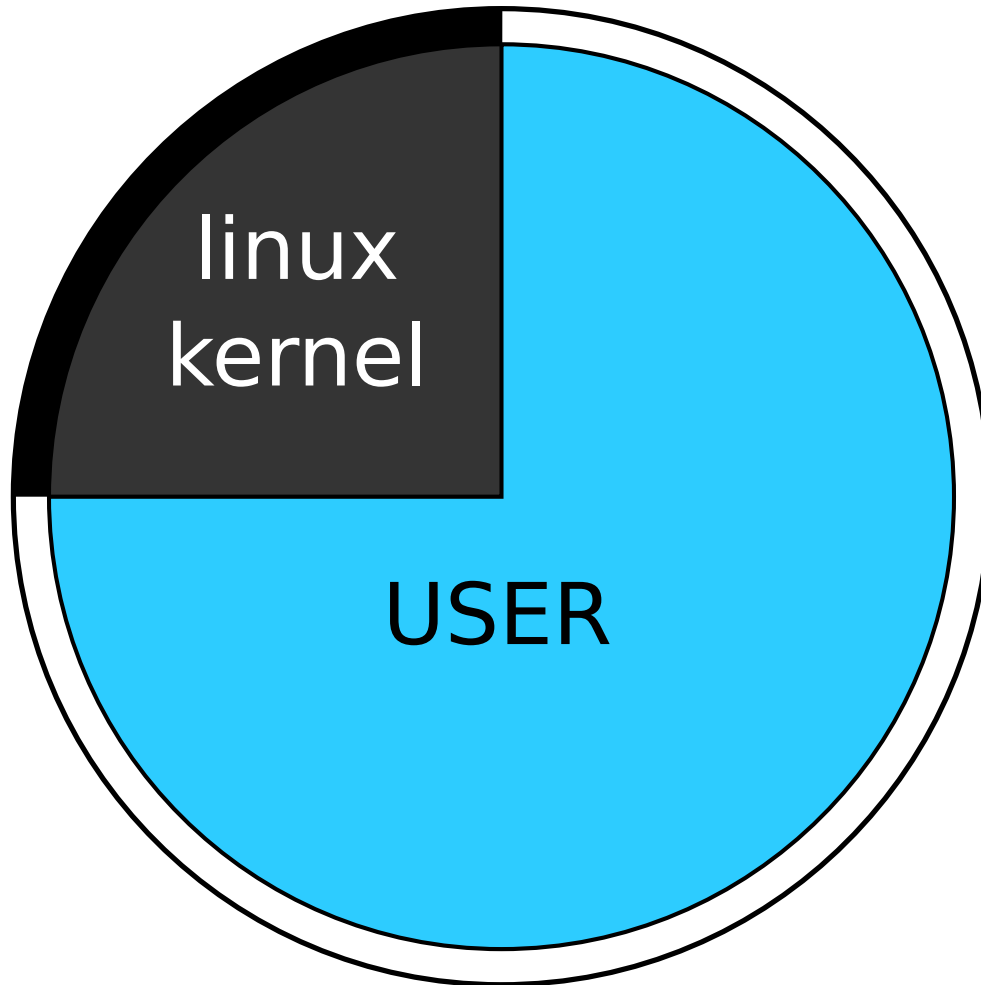


memory layout

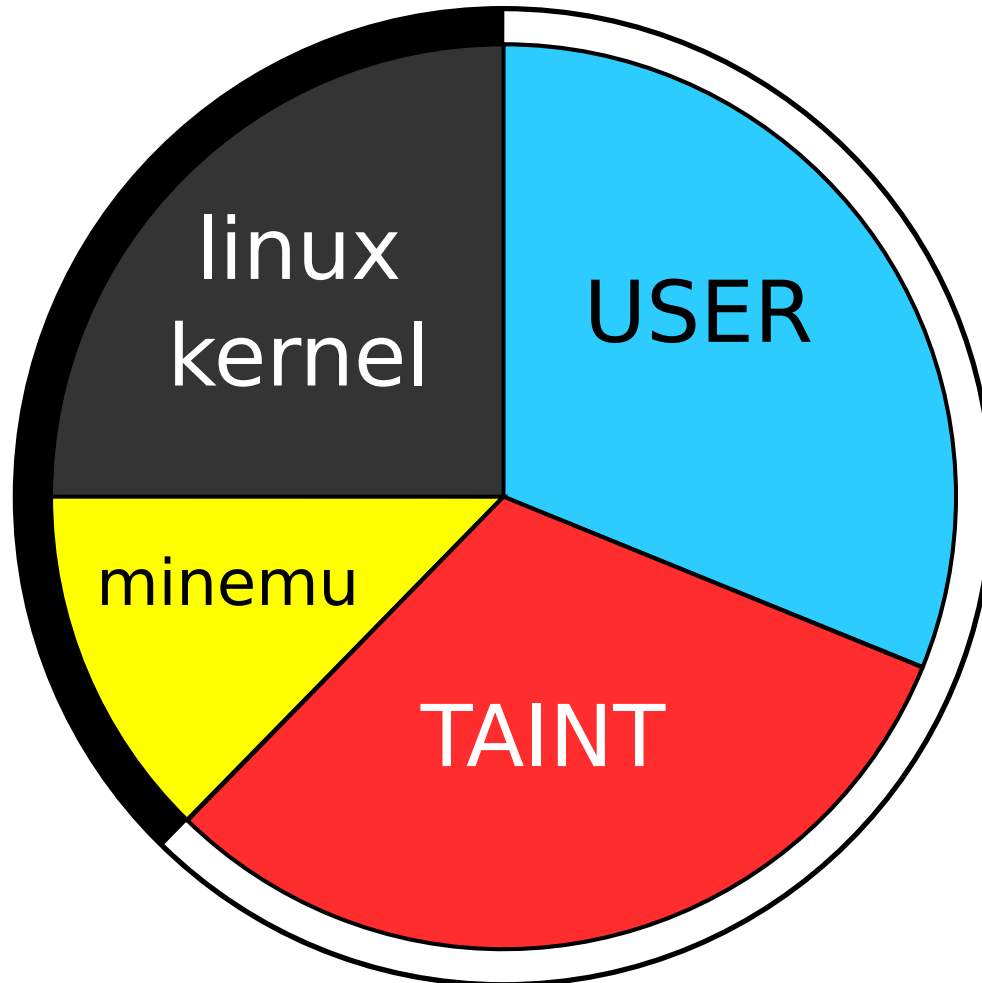
use SSE registers to hold taint



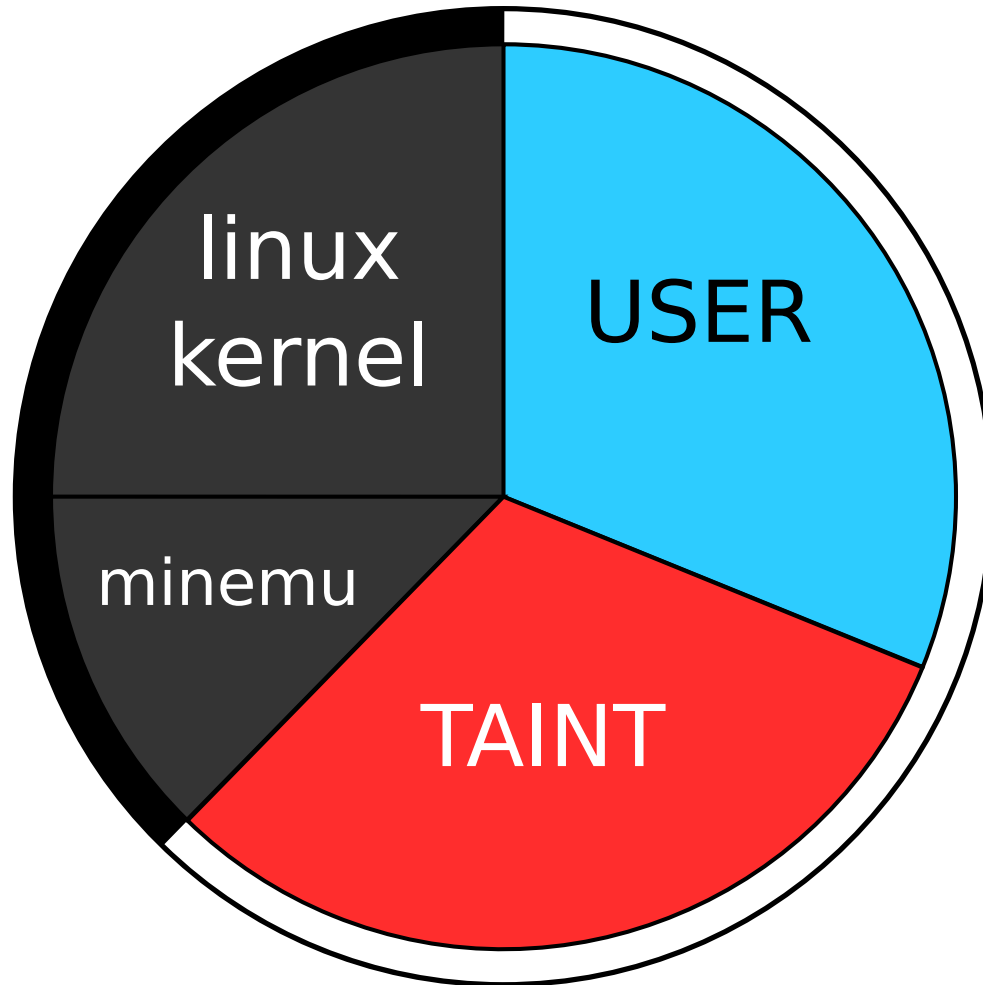
Memory layout (linux)



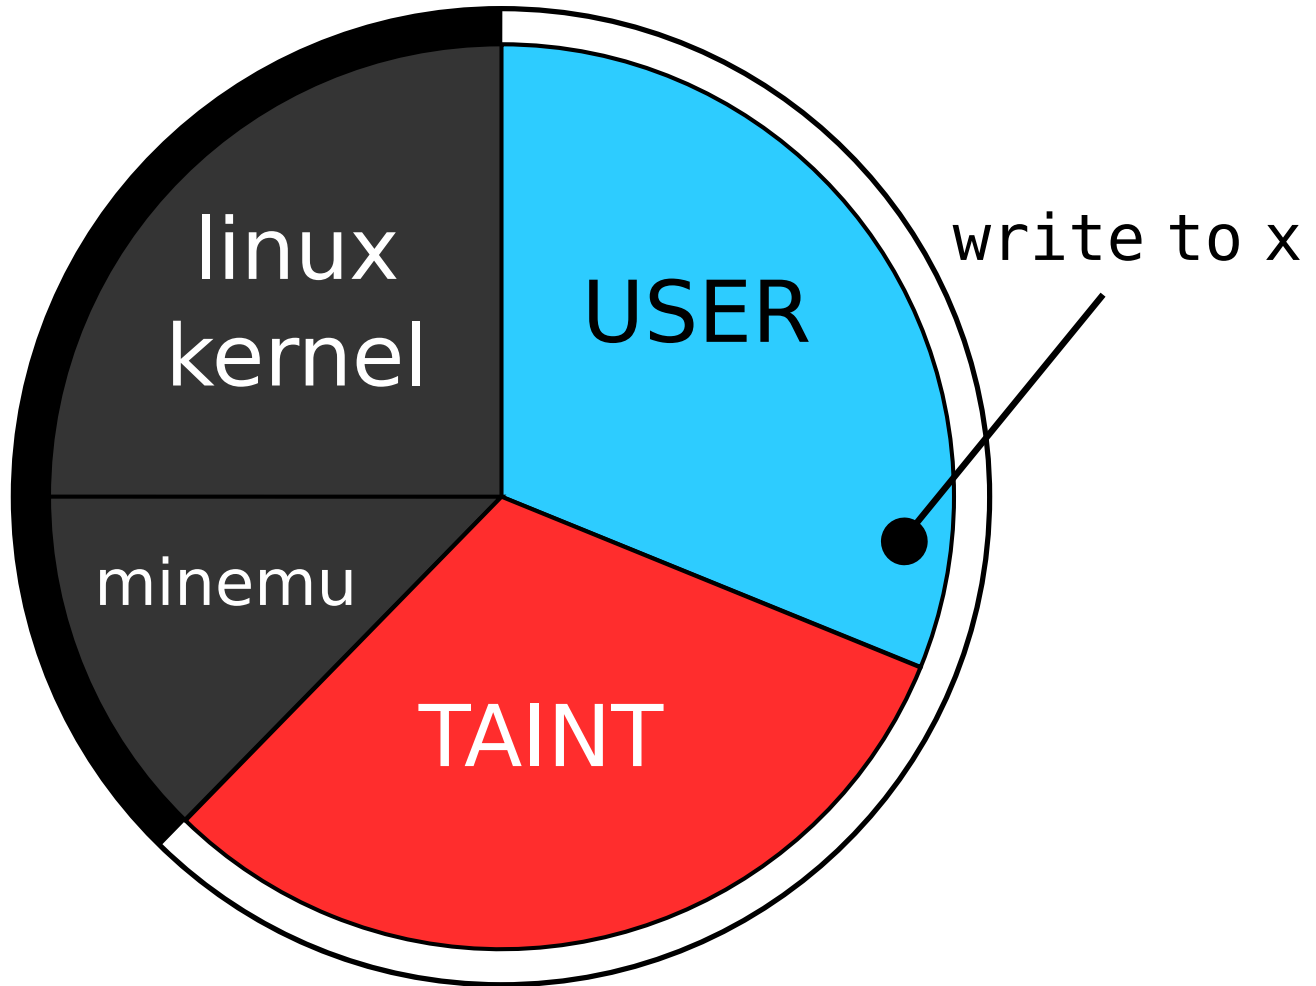
Memory layout (minemu)



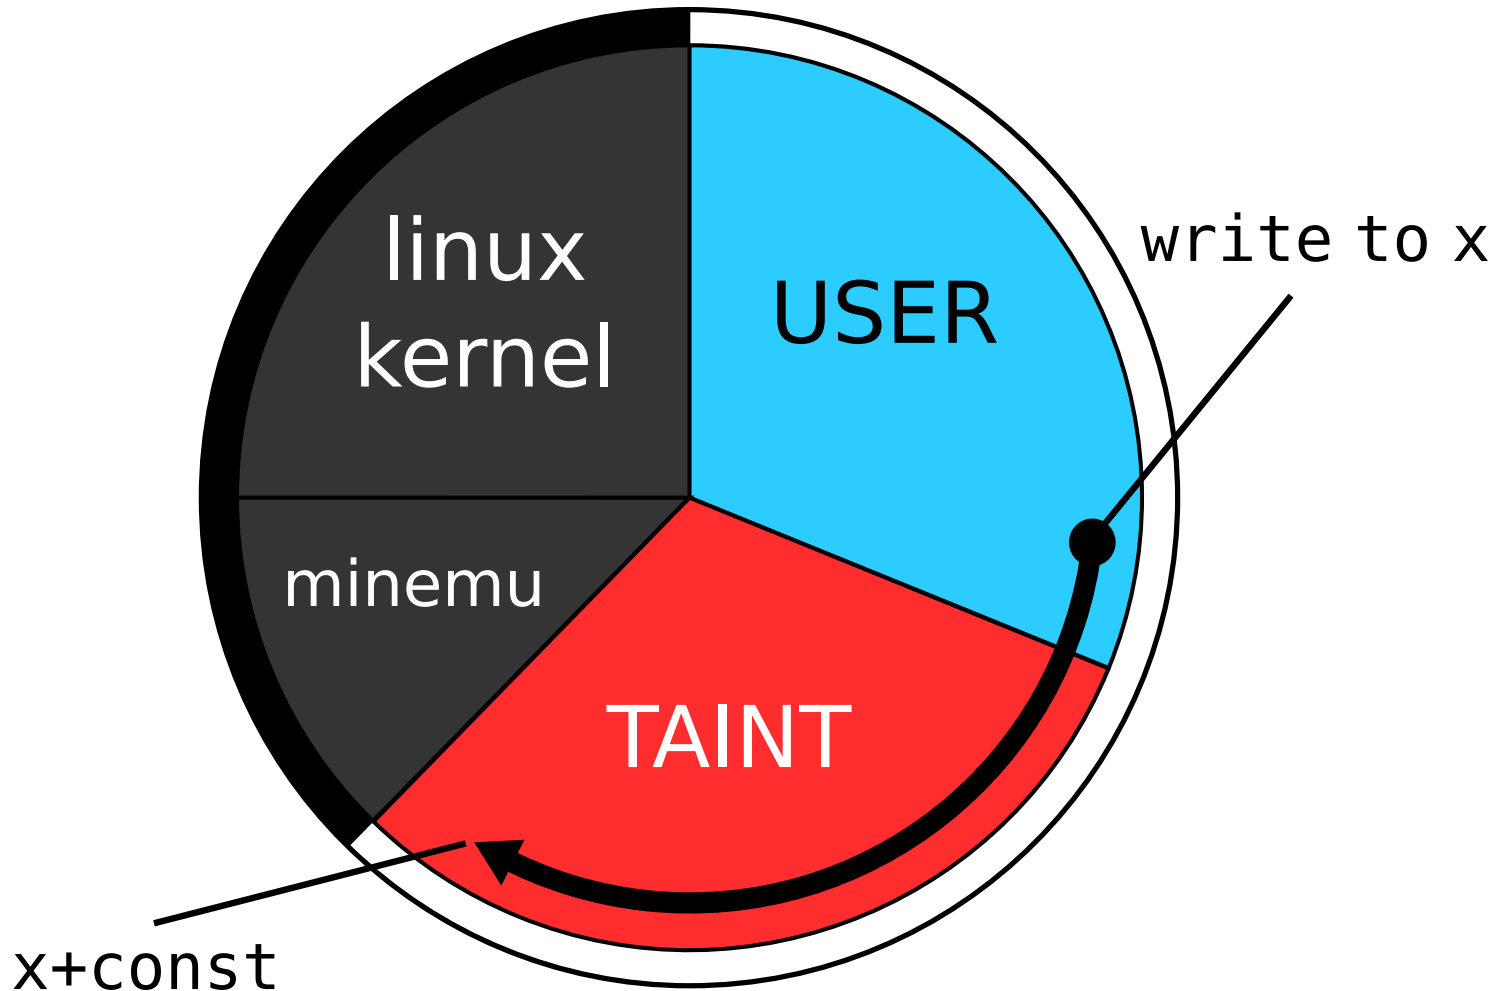
Memory layout (minemu)



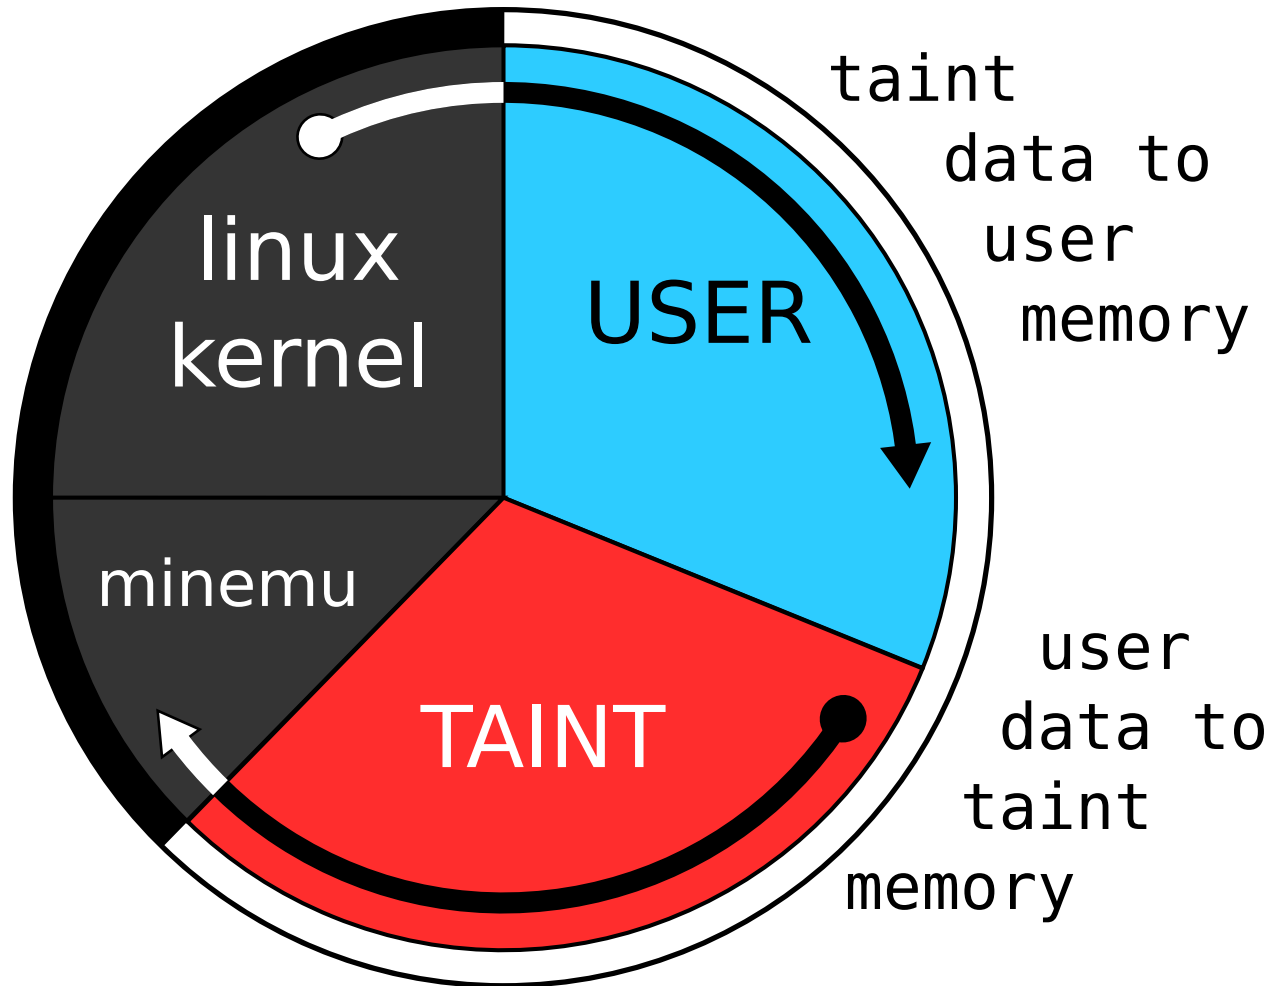
Memory layout (minemu)



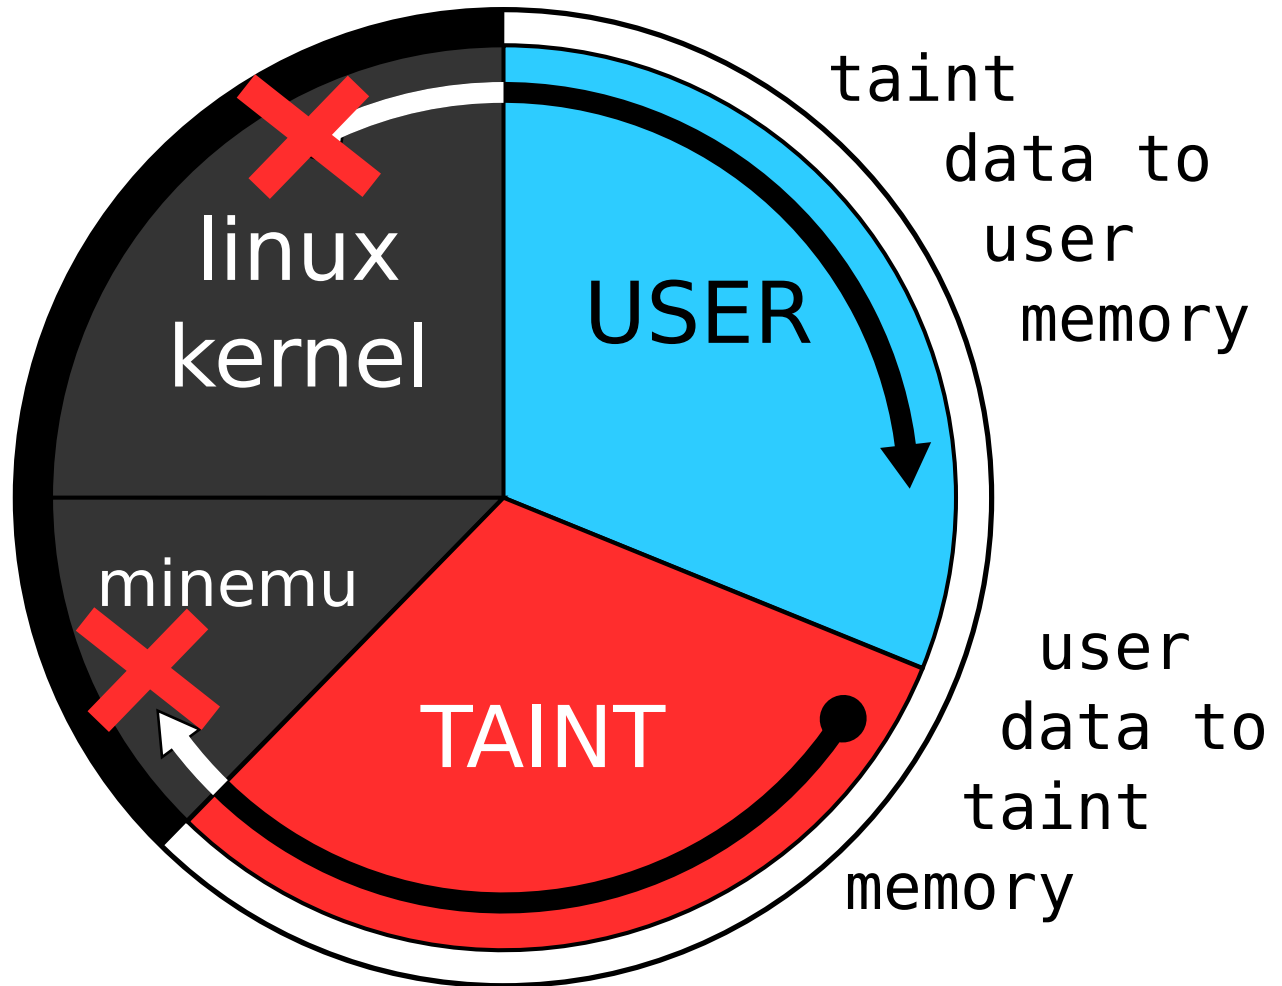
Memory layout (minemu)



Memory layout (minemu)



Memory layout (minemu)



Addressing shadow memory

```
mov EAX, (EDX)
```

Addressing shadow memory

```
mov EAX, (EDX)
```

address:

EDX

Addressing shadow memory

```
mov EAX, (EDX)
```

address:

EDX

taint:

EDX+**const**

Addressing shadow memory

```
mov EAX, (EDX+EBX*4)
```

Addressing shadow memory

```
mov EAX, (EDX+EBX*4)
```

address:

$EDX+EBX*4$

Addressing shadow memory

```
mov EAX, (EDX+EBX*4)
```

address:

$EDX+EBX*4$

taint:

$EDX+EBX*4+const$

Addressing shadow memory

```
push ESI
```

Addressing shadow memory

push ESI

address:

ESP

Addressing shadow memory

push ESI

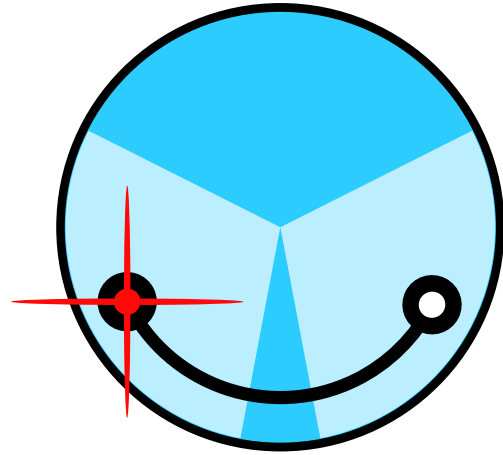
address:

ESP

taint:

ESP+**const**

Is this slowness fundamental?

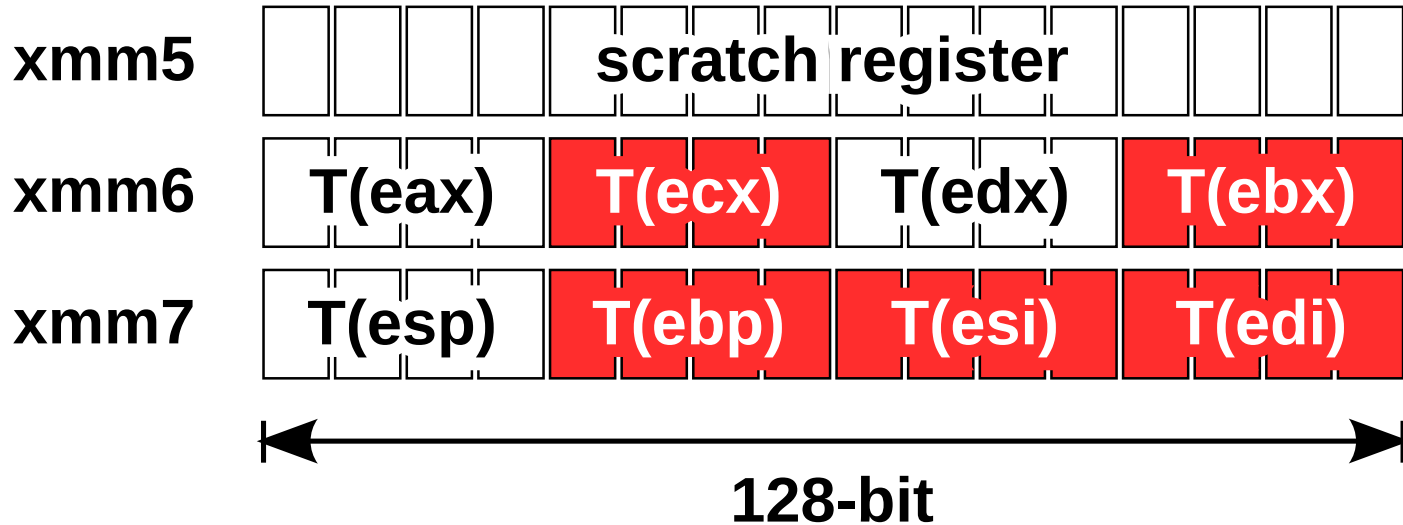


minemu

fast emulator
memory layout

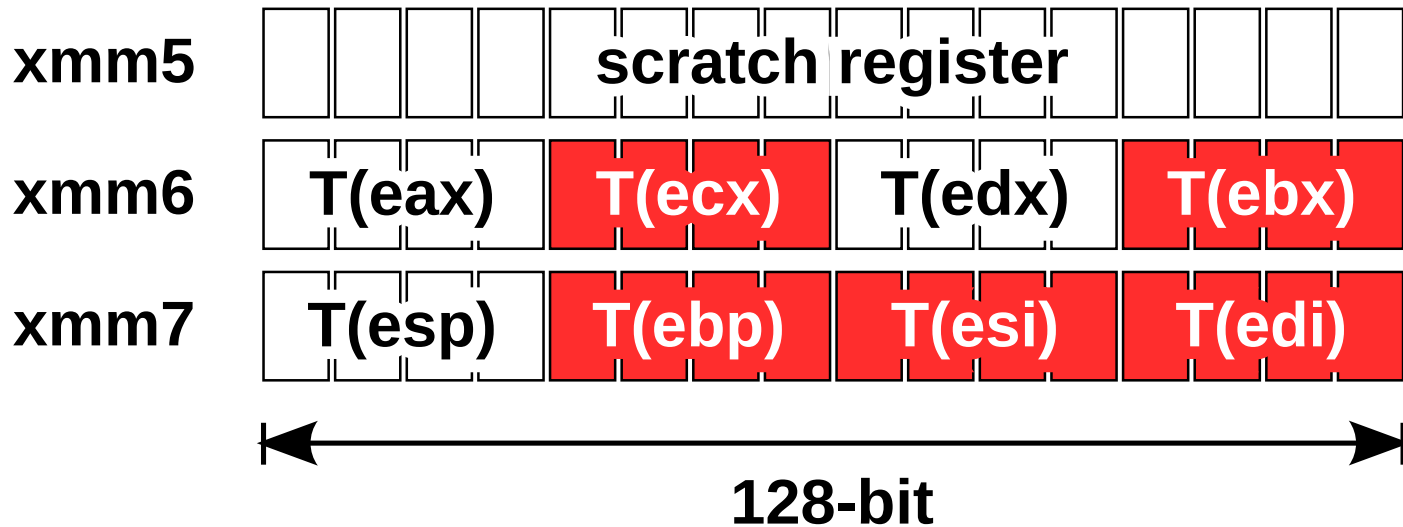
► use SSE registers to hold taint

Taint propagation in SSE registers



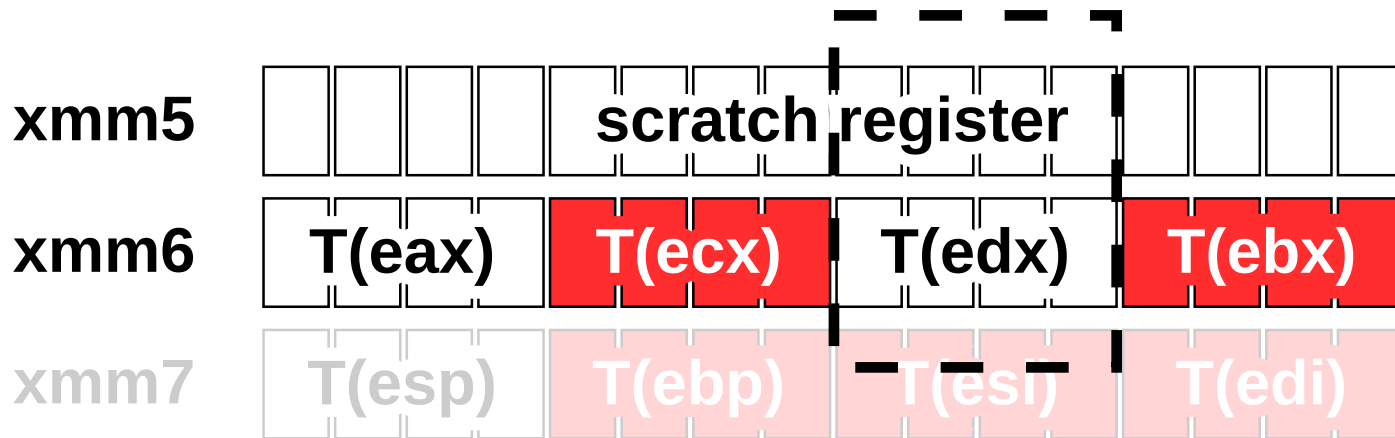
Taint propagation in SSE registers

add EDX, x



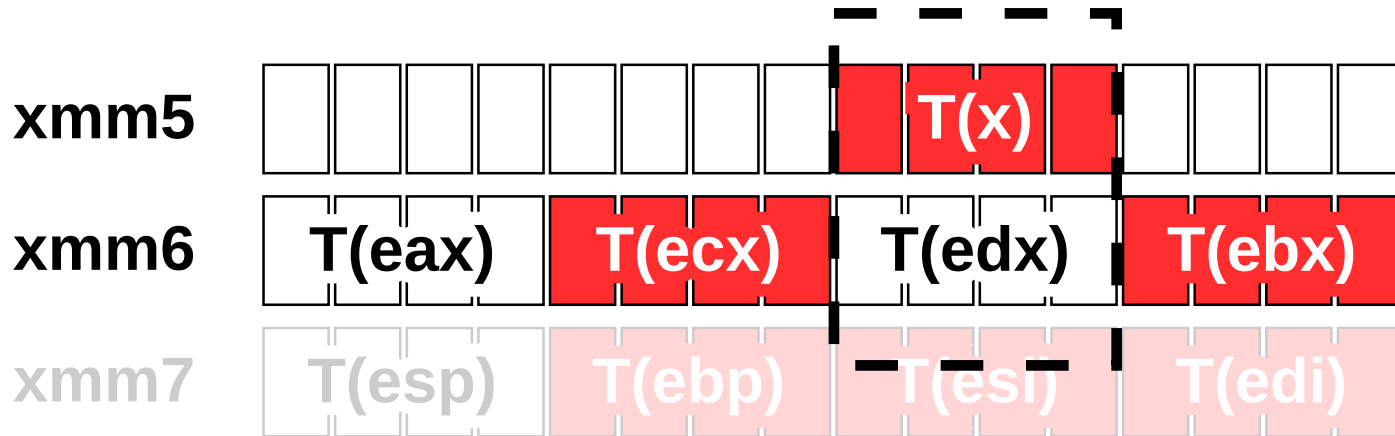
Taint propagation in SSE registers

add EDX, x



Taint propagation in SSE registers

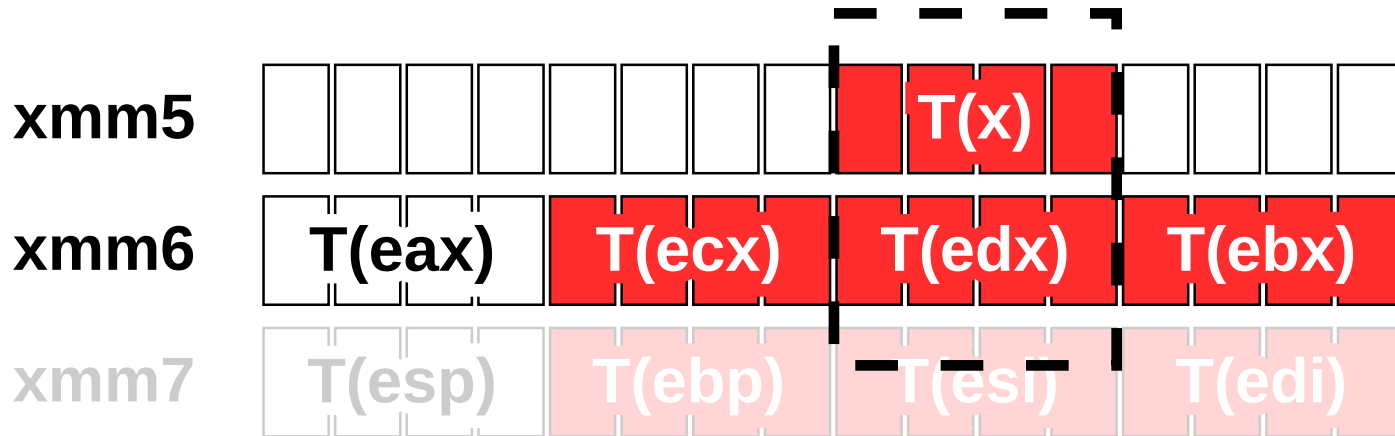
add EDX, x



vector insert

Taint propagation in SSE registers

add EDX, x



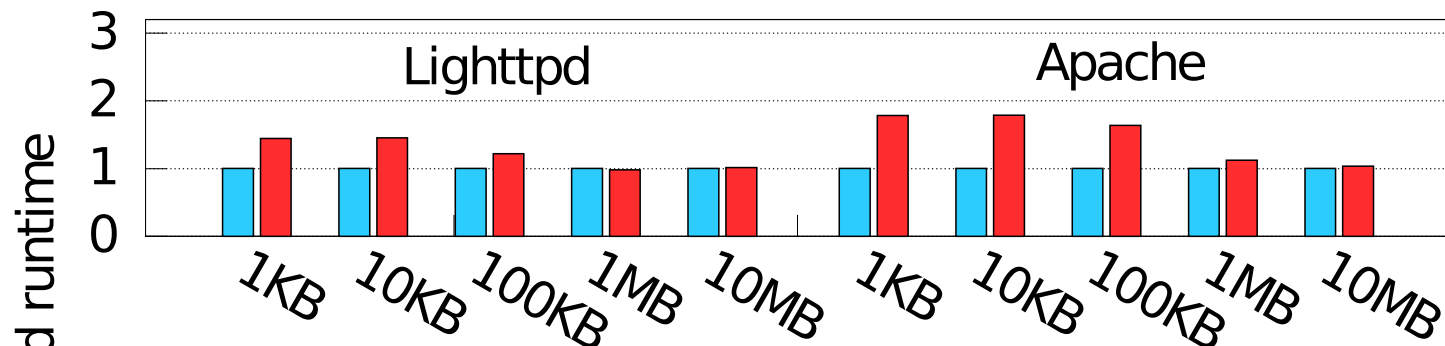
or

Effectiveness

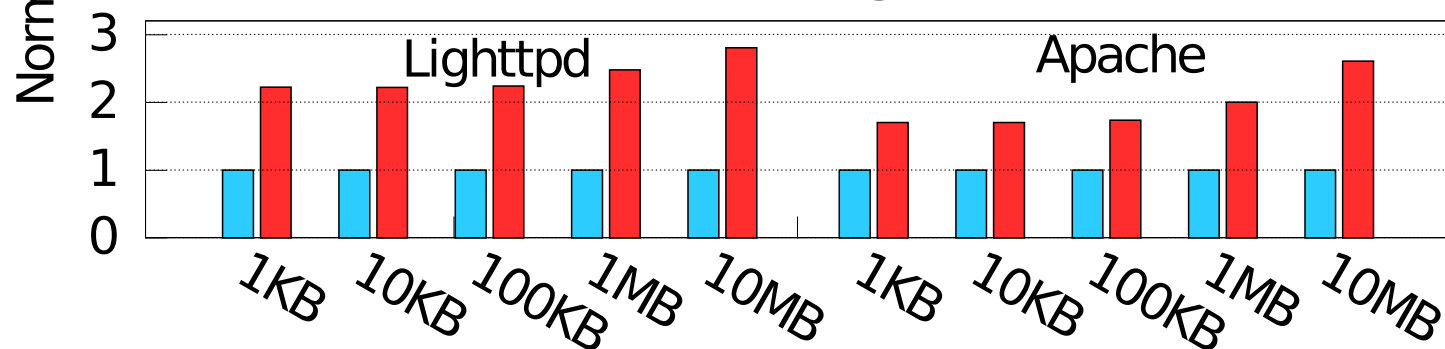
Application	Type of vulnerability	Security advisory
Snort 2.4.0	Stack overflow	CVE-2005-3252
Cyrus imapd 2.3.2	Stack overflow	CVE-2006-2502
Samba 3.0.22	Heap overflow	CVE-2007-2446
Memcached 1.1.12	Heap overflow	CVE-2009-2415
Nginx 0.6.32	Buffer underrun	CVE-2009-2629
Proftpd 1.3.3a	Stack overflow	CVE-2010-4221
Samba 3.2.5	Heap overflow	CVE-2010-2063
Telnetd 1.6	Heap overflow	CVE-2011-4862
Ncompress 4.2.4	Stack overflow	CVE-2001-1413
Iwconfig V.26	Stack overflow	CVE-2003-0947
Aspell 0.50.5	Stack overflow	CVE-2004-0548
Htget 0.93	Stack overflow	CVE-2004-0852
Socat 1.4	Format string	CVE-2004-1484
Aeon 0.2a	Stack overflow	CVE-2005-1019
Exim 4.41	Stack overflow	EDB-ID#796
Htget 0.93	Stack overflow	
Tipxd 1.1.1	Format string	OSVDB-ID#12346

Performance

HTTP

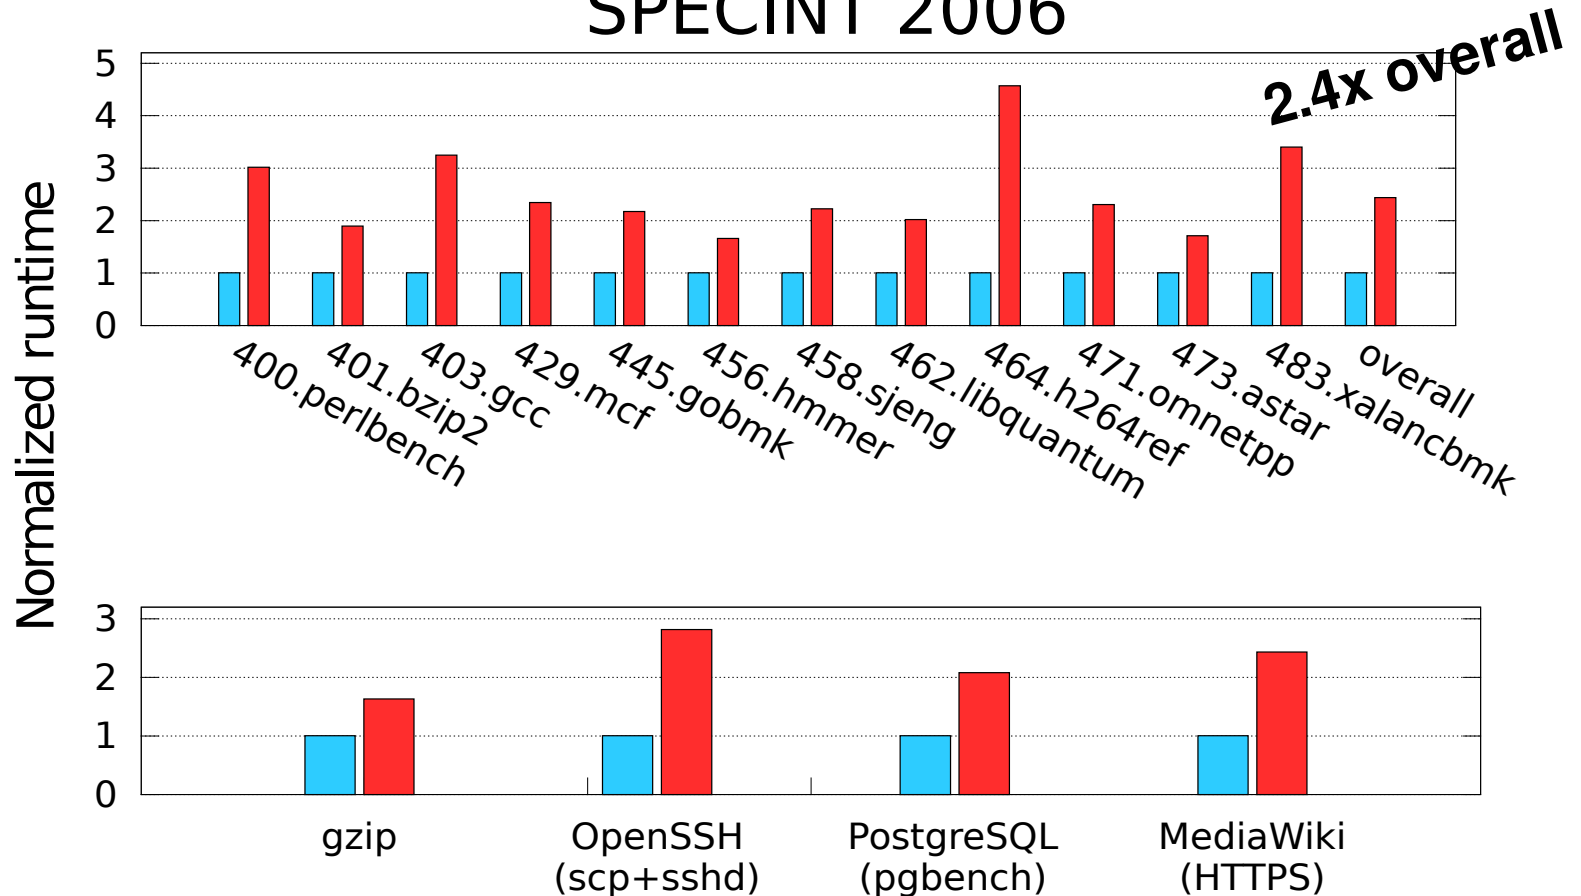


HTTPS



Performance

SPECINT 2006



Limitations

Limitations

Doesn't prevent memory corruption, only acts when the untrusted data is used for arbitrary code execution.

Limitations

Tainted pointer dereferences

```
tainted_pointer->some_field = useful_untainted_value;
```

Limitations

Tainted pointer dereferences

```
tainted_pointer->some_field = useful_untainted_value;
```

propagation can lead to false positives:

```
dispatch_table[checked_input]();
```

Limitations

Taint whitewashing

```
out = latin1_to_ascii[in];
```

Limitations

Format string attacks:

```
printf("%65534s %123$hn"); // Propagates taint in glibc
```

```
printf("FillerFiller...%123$hn"); // Does not :-)
```

Limitations

Does not protect against non-control-flow exploits

Limitations

Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```

Limitations

Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```


Limitations

Does not protect against non-control-flow exploits

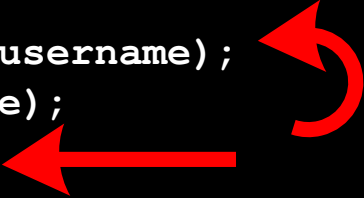
```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```



Limitations

Does not protect against non-control-flow exploits

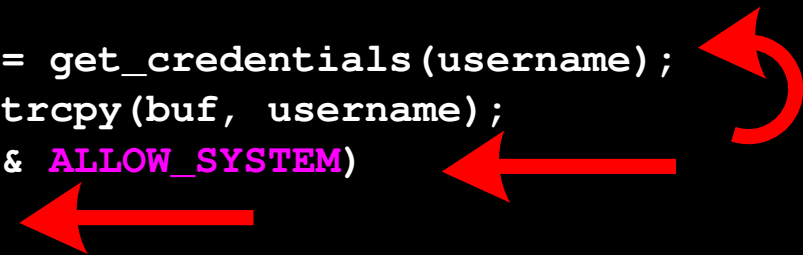
```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```

Two red arrows are present. One is a curved arrow pointing from the right towards the 'strcpy(buf, username);' line. The other is a straight arrow pointing from the right towards the 'ALLOW_SYSTEM' constant in the 'if' statement.

Limitations

Does not protect against non-control-flow exploits

```
void try_system(char *username, char *cmd)
{
    int user_rights = get_credentials(username);
    char buf[16] ; strcpy(buf, username);
    if (user_rights & ALLOW_SYSTEM)
        system(cmd);
    else
        log_error("user %s attempted login", buf);
}
```



The diagram illustrates a control flow jump. A red arrow points from the 'if' statement to the 'else' block, indicating that the program execution will skip the 'system(cmd);' call and proceed to the 'log_error' call if the condition is false. Another red arrow points from the 'system(cmd);' line back to the 'if' statement, indicating the flow continues to the next line after the 'if' block if the condition is true.

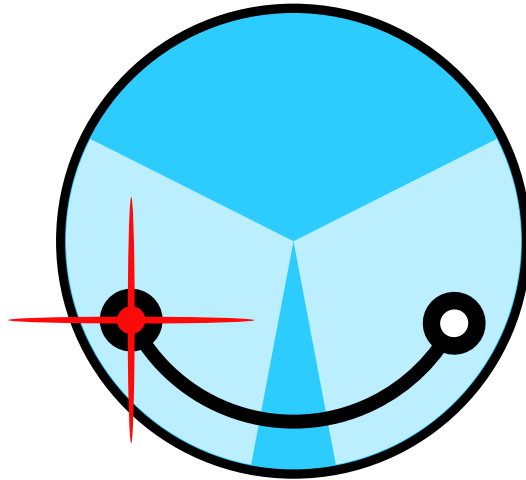
in some cases we can add validation hooks.

`mysql_query()` can be hooked to check for taint outside of literals in SQL queries.

in some cases we can add validation hooks.

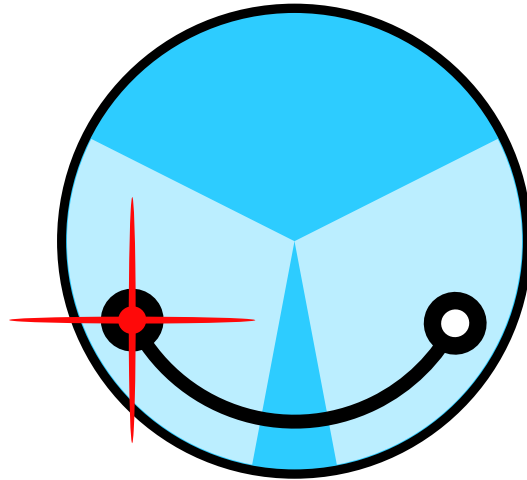
`mysql_query()` can be hooked to check for taint outside of literals in SQL queries.

`_IO_vfprintf()` in glibc can be hooked to check format strings for taint.



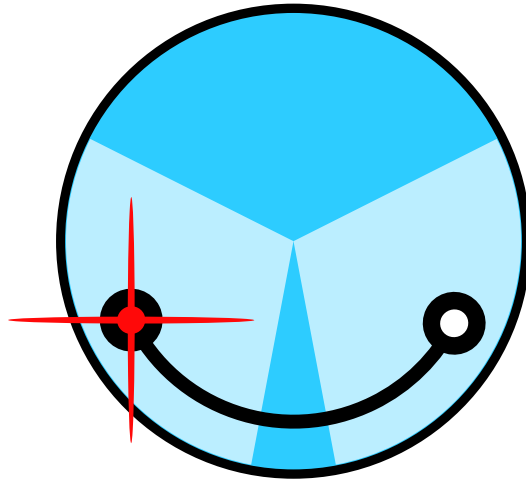
Demo

```
demo@demo:~# ./minemu bash
```



Minemu

```
git clone https://minemu.org/code/minemu.git
```



Minemu

```
git clone https://minemu.org/code/minemu.git
```

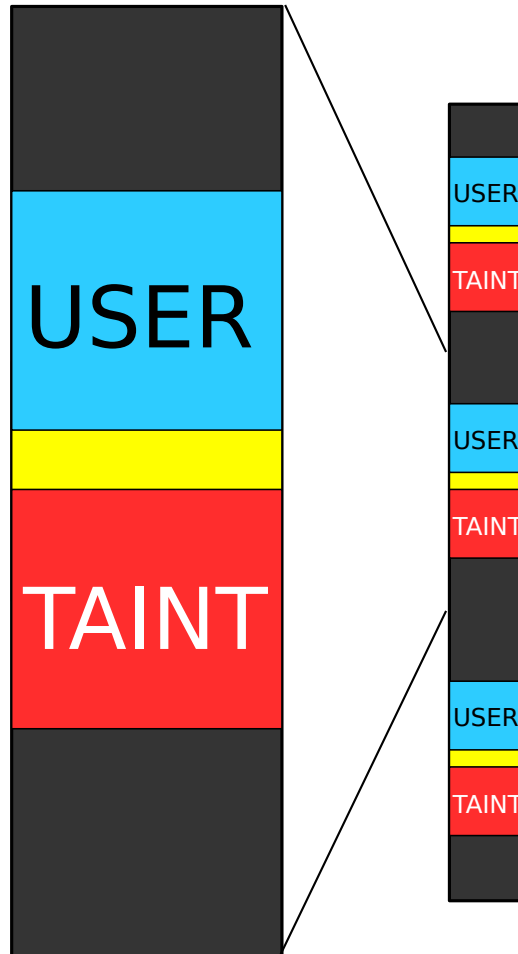
any questions?

<https://minemu.org/vms/>

5f1ee00029e2c68699a7670de7aef02e	minemu-demo.ova
c4ee74155a858676bfb54e1fcfb6db0e	minemu-demo.qcow2
5b8b910c38901f43d406a21fe9767822	minemu-demo.vdi.gz
7ba81ae9d35bfa05a70068a804a331ac	minemu-demo.vmdk.gz
c37acdc455ebac700139f60da621bc38	minemu-demo.xml

minemu needs CPU with SSE 4.1

Memory layout (64 bit)



Memory layout (64 bit) alternative

